Minor in AI
Convolution Neural Network

January 2, 2025
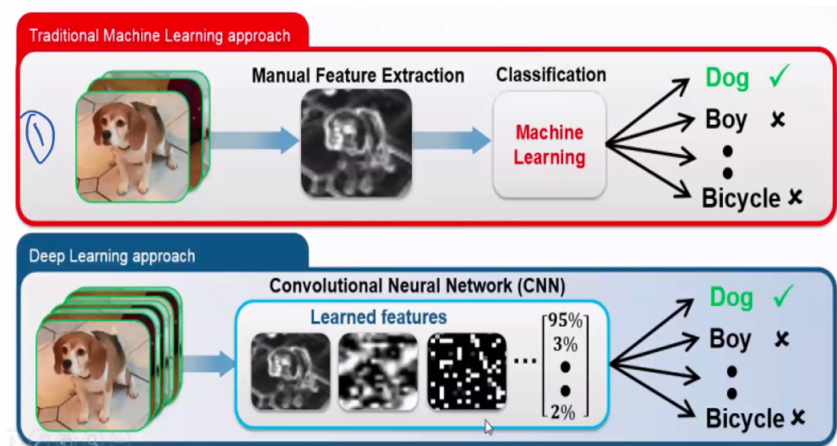
Figure 1: Deep Learning vs Machine Learning

1   Introduction

Convolutional Neural Networks (CNNs) are a class of deep neural networks designed specifically for tasks involving structured data such as images, video sequences, and time-series data. Their architecture mimics the connectivity pattern of neurons in the human brain, where small clusters of neurons respond to stimuli in their receptive field.

CNNs are widely used in:

- Image and video recognition

- Object detection

- Medical image analysis

- Natural language processing (e.g., sentence classification)

2   Traditional Machine Learning vs Deep Learning

- **Feature Extraction:**

  - Traditional ML relies on manually designed features.
  - Deep Learning automatically extracts hierarchical features.

- **Scalability:**

  - Traditional ML struggles with large-scale data.
  - Deep Learning thrives on massive datasets using GPUs/TPUs.

- **Performance:**

  - Traditional ML models like SVM and Random Forests perform well on small datasets.
  - Deep Learning surpasses in tasks involving unstructured data such as images and videos.

Figure 2: Convolution in 2D

3 Steps in a CNN

1. **Convolution:** Applies filters (kernels) to extract features such as edges, textures, or complex patterns.

2. **ReLU:** Introduces non-linearity by applying $f(x) = \max(0, x)$, enabling the model to learn non-linear decision boundaries.

3. **Pooling:** Reduces dimensionality while retaining essential features, improving computational efficiency.

4. **Fully Connected Layer:** Combines extracted features for final classification or regression output.

4 Convolution in 2D

Convolution is a mathematical operation that processes an input matrix using a smaller matrix called a kernel (or filter). It involves sliding the kernel over the input matrix and computing the weighted sum of the overlapping elements. The result is stored in an output matrix. To understand the formula:

1. For each position (i, j) in the output matrix Y, place the kernel K over the input matrix X, aligning the top-left corner of the kernel with the position (i, j) in X.

2. Multiply each element of the kernel $K[m, n]$ by the corresponding element in the input matrix $X[i + m, j + n]$.

3. Sum all these multiplied values to compute a single number, which becomes the value of $Y[i, j]$.

4. Slide the kernel to the next position and repeat the process until all positions in Y are filled.

In simpler terms, convolution applies a filter (kernel) to the input to detect patterns, reduce dimensionality, or create new representations of the data. Mathematically, it can be written as:

$$Y[i,j] = \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} X[i+m, j+n] \cdot K[m,n]$$

Where:

- X: The input matrix, which represents the data to be processed (e.g., an image or feature map).

- K: The kernel or filter, a small matrix containing weights used to extract specific patterns from the input.

- Y: The output matrix, which stores the results of the convolution operation.

- i, j: The row and column indices in the output matrix Y.

- k: The size (height and width) of the kernel K, assuming it is square.

- m, n: Indices that traverse the rows and columns of the kernel.

4.1   Example: Matrix Convolution (Single Step Calculation)

Let us perform a single step of the convolution operation using the following matrices:

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} \cdots \end{bmatrix}$$

$$\text{Input Matrix (Slice)} = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}, \quad \text{Kernel} = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

The convolution operation involves element-wise multiplication followed by summation:

Step-by-step Calculation:

$$(1 \cdot 1) + (1 \cdot 0) + (1 \cdot 1)$$
$$+(0 \cdot 0) + (1 \cdot 1) + (1 \cdot 0)$$
$$+(0 \cdot 1) + (0 \cdot 0) + (1 \cdot 1)$$

$$\text{Result} = 1 + 0 + 1 + 0 + 1 + 0 + 0 + 0 + 1 = 4$$

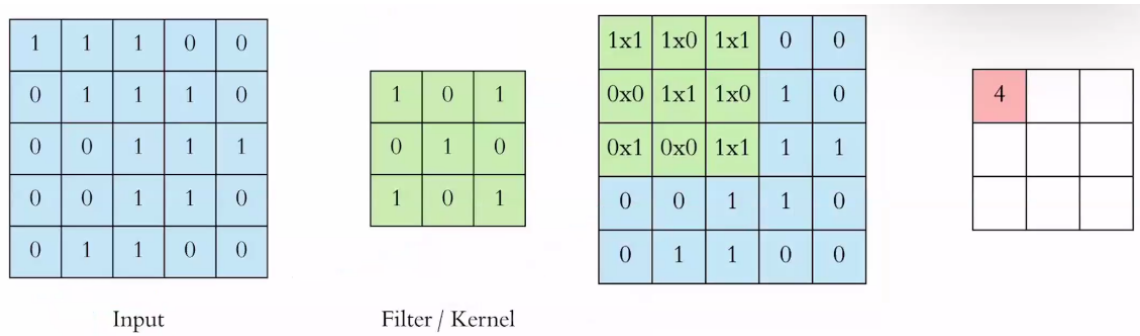Thus, the value of the convolution operation at this position is: 4

Figure 3: Convolution Operation

4.2 Padding

Padding ensures the output size is preserved or adjusted based on the application:

- **Valid Padding:** No padding is applied, reducing output size.

- **Same Padding:** Adds zeros around the input to maintain the output size.

4.3 Stride

Stride refers to the step size with which the kernel moves. It impacts:

- **Output Size:** Larger strides result in smaller output dimensions.

- **Computational Cost:** Higher strides reduce computation at the cost of detail loss.

5 Activation Functions

Activation functions introduce non-linearity into the model. Common types:

- **Sigmoid:** Squashes input into $(0, 1)$. Useful for probabilities.
$$f(x) = \frac{1}{1 + e^{-x}}$$

- **Tanh:** Squashes input into $(-1, 1)$. Centered at zero.
$$f(x) = \tanh(x)$$

- **ReLU:** Replaces negative values with zero.
$$f(x) = \max(0, x)$$

- **Leaky ReLU:** Allows a small gradient for negative values.
$$f(x) = \begin{cases} x & x \geq 0 \\ \alpha x & x < 0 \end{cases}$$

- **ELU:** Exponential Linear Unit, smoothes ReLU's zero gradient issue.

Figure 4: Activation Functions

6 Region of Interest (ROI)

A **Region of Interest (ROI)** refers to a specific part of an image or data that is particularly important for the problem at hand. Instead of processing the entire image, focusing on the ROI allows us to work on areas that hold meaningful information, reducing computational complexity and improving efficiency.

6.1 Why ROI is Important

- **Efficiency:** By processing only the ROI, we save resources and time, especially for high-resolution images.

- **Focus on Relevant Data:** Not all parts of an image are equally important. For example:
    - In medical imaging, the ROI could be a tumor in an X-ray or MRI scan.
    - In facial recognition, the ROI is the face, excluding the background.

- **Improved Accuracy:** Concentrating on the ROI reduces noise and irrelevant data, which can confuse the model.

6.2 Applications of ROI

- **Medical Imaging:** Detecting abnormalities such as tumors or fractures.

- **Object Detection:** Highlighting and identifying objects like cars, pedestrians, or animals in an image.

- **Image Segmentation:** Dividing an image into meaningful sections, where each section has a specific purpose.

7 Limitations of 2D Convolution

While 2D convolution is highly effective, it comes with certain limitations:

7.1 Fixed Receptive Field

- The receptive field is the area of the input image covered by the kernel at a given time.

- A fixed kernel size may miss patterns that span a larger area, such as global textures or large-scale structures.

7.2 Loss of Contextual Information

- Convolution operations focus on local patterns, such as edges or small textures.

- As a result, they may not capture relationships between distant parts of an image, which is essential for understanding the overall structure.

7.3 High Computational Cost

- For high-resolution images, the number of computations increases significantly due to the large input size.

- This can lead to slower processing and the need for more powerful hardware.

7.4 Limited Temporal Understanding

- 2D convolution cannot model temporal changes in data, such as variations across video frames or sequences of events.

- For such cases, specialized architectures like 3D Convolutions or Recurrent Neural Networks (RNNs) are required.

7.5 Boundary Effects

- Convolution may produce distorted results at the boundaries of an image due to incomplete coverage of the kernel.

- Padding partially mitigates this issue, but it introduces artificial data (zeros or other values) into the computation.

7.6 Sensitivity to Rotation and Scaling

- Convolutional filters are sensitive to the orientation and size of patterns in the image.

- Features such as edges may not be detected properly if the object is rotated or scaled.

8 Conclusion

Convolutional Neural Networks (CNNs) are a cornerstone of modern deep learning, excelling in tasks involving structured data such as images and videos. Their architecture, inspired by the human brain's visual processing, allows them to learn hierarchical features from raw data without manual feature engineering.

8.1 Key Takeaways

- **Loss Functions:** The loss function is a critical component for supervised learning, quantifying the difference between predicted outputs and ground truth. Examples include Mean Squared Error (MSE) for regression and Cross-Entropy Loss for classification.

- **CNN Workflow:** The steps of a CNN—Convolution, ReLU, Pooling, and Fully Connected layers—together enable efficient feature extraction and decision-making. Each step plays a distinct role:

  - Convolution captures local patterns using kernels.
  - ReLU introduces non-linearity, enabling the model to learn complex patterns.
  - Pooling reduces dimensionality, improving computational efficiency.
  - Fully connected layers integrate features for final predictions.

- **Region of Interest (ROI):** Focusing on specific regions of interest improves model efficiency and accuracy by reducing noise and processing only the most relevant parts of an image. Applications include medical imaging, object detection, and segmentation.

- **Activation Functions:** Non-linear activation functions such as Sigmoid, Tanh, ReLU, and Leaky ReLU allow neural networks to model complex relationships in data. Each activation function has specific use cases based on the nature of the task and data.

- **Limitations of 2D Convolution:** While 2D convolutions are powerful, they face challenges such as:

  - Fixed receptive fields, limiting their ability to capture global patterns.
  - Loss of contextual information from distant parts of an image.
  - High computational cost for large-scale data.
  - Inability to model temporal relationships, as seen in sequential data.
  - Sensitivity to changes in rotation, scale, and orientation.

  Techniques like 3D convolutions, padding, and architectural innovations aim to address these limitations.

- **Comparison with Traditional Machine Learning:** CNNs outperform traditional approaches in tasks involving unstructured data by learning features directly from raw data. They also scale better with larger datasets and complex tasks.

9 Appendix

9.1 Loss Functions

In supervised learning, the **loss function** serves as the guiding metric for model optimization. It calculates the error between the predicted output (y_{pred}) and the ground truth (y_{true}).

Examples:

- **Mean Squared Error (MSE):** Common for regression tasks.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^{n} \left(y_{\text{true}}^{(i)} - y_{\text{pred}}^{(i)} \right)^2$$

- **Cross-Entropy Loss:** Common for classification tasks.

$$\text{Cross-Entropy} = -\frac{1}{n} \sum_{i=1}^{n} \sum_{c=1}^{C} y_{\text{true},c}^{(i)} \log(y_{\text{pred},c}^{(i)})$$